




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Continuous validation of software performance at CSCS: How can you contribute?

User Lab Day 2022

Victor Holanda Rusu, CSCS

September 2nd, 2022

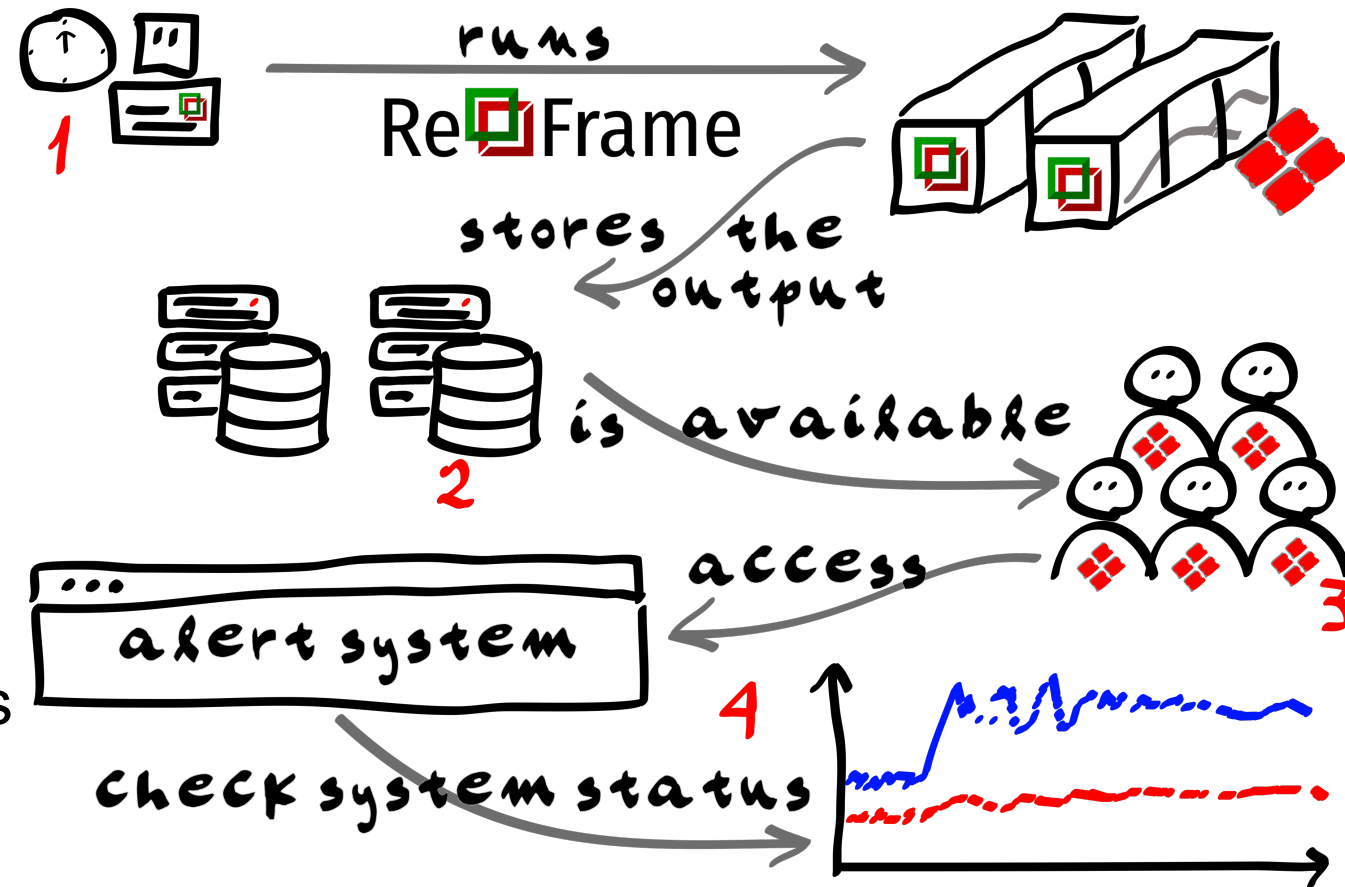


Regression testing at CSCS

This is all about quality of service

■ Everyday at CSCS...

1. A robot runs a **selection of ReFrame checks** on **every** User Lab system
2. The **test results** are **stored** in a dedicated database
3. CSCS **engineers** check the results of the tests
4. And control if there is any **performance degradation** in any system






Regression testing at CSCS

How do the results look like?

Results improved upon system change

But no significant changes happened for other metrics


Maybe your application behaves like this one?



Regression testing at CSCS

Where can you contribute?

- For us it is very important to monitor the system from the user perspective
- On top of system testing and synthetic micro benchmarks
- We also test
 - Tools that users can run
 - Applications and libraries that users can run
- But the day only has 24 hours and we are not experts in every single application running on our systems





Contributing to ReFrame

What is there for you?

- Help us identify changes to the **current system** that affect your workflow
- Help us identify changes to the **next system** that affect your workflow
- **Reuse** your tests in **your own lab** environment
- Tests will be **portable** across different HPC sites
- We will be able to provide **better** overall quality of service **to you**
- Tests **can be re-used** for the next project proposal **technical report**

Help us, help you!

Join the ReFrame team





ReFrame at CSCS



Tests pipeline

What is ReFrame?

- ReFrame is an Open Source Regression Testing framework that started at CSCS and has evolved to a community driven project
- It was designed from the ground up for writing portable validation, regression and performance tests running from laptops to Top500 supercomputers
- First public release in May 2017 on Github
 - 42 contributors, 78 forks, 154 stars
 - 70 releases
- Contains several different characteristics, among them
 - Implements a Python eDSL allowing to write tests in a high-level declarative way
 - Tests are composable by design and can be extended or reused across sites
 - Multi-dimensional test parameterisation
 - Support for native and containerised runs
 - Seamless integration with Gitlab CI/CD pipelines

ReFrame – Efficient System and Application Performance Testing

ReFrame tests

What do they look like?

lammps_check.py

```
6 import os
7
8 import reframe as rfm
9 import reframe.utility.sanity as sn
10
11 Tests are Python Classes
12 class LAMMPSCheck(rfm.RunOnlyRegressionTest):
13     scale = parameter(['small', 'large'])
14     modules = ['cray-python', 'LAMMPS']
15     tags = {'external-resources', 'maintenance', 'production'}
16     maintainers = ['LM']
17     strict_check = False
18     extra_resources = {
19         'switches': {
20             'num_switches': 1
21         }
22     }
```

Parameterise
the tests

Define what's the
performance metric

Define what's the
sanity metric

Make use of modules
available at CSCS

We can store the input files for
you in case they are very large

```
24 @run_after('init')
25 def setup_by_system(self):
26     # Reset sources dir relative to the SCS apps prefix
27     self.sourcesdir = os.path.join(self.current_system.resourcesdir,
28                                     'LAMMPS')
29     if self.current_system.name in ['eiger', 'pilatus']:
30         self.valid_prog_environs = ['cpeGNU']
31     else:
32         self.valid_prog_environs = ['builtin']
33
34 @performance_function('timesteps/s')
35 def perf(self):
36     return sn.extractsingle(r'\s+(?P<perf>\S+) timesteps/s',
37                             self.stdout, 'perf', float)
38
39 @sanity_function
40 def assert_energy_diff(self):
41     energy_reference = -4.6195
42     energy = sn.extractsingle(
43         r'\s+500000(\s+\S+){3}\s+(?P<energy>\S+)\s+\S+\s\n',
44         self.stdout, 'energy', float)
45     energy_diff = sn.abs(energy - energy_reference)
46     return sn.all([
47         sn.assert_found(r'Total wall time:', self.stdout),
48         sn.assert_lt(energy_diff, 6e-4)
49     ])
```

ReFrame tests

What do they look like?

lammps_check.py

Set the
executable
and its
options

Anotate to make tests runnable

```
52 @rfm.simple_test
53 class LAMMPSGPUCheck(LAMMPSCheck):
54     valid_systems = ['daint:gpu']
55     executable = 'lmp_mpi'
56     executable_opts = ['-sf gpu', '-pk gpu 1', '-in in.lj.gpu']
57     variables = {'CRAY_CUDA_MPS': '1'}
58     num_gpus_per_node = 1
59     refs_by_scale = {
60         'small': {
61             'dom:gpu': {'perf': (3456.792, -0.10, None, 'timesteps/s')},
62             'daint:gpu': {'perf': (1566.979, -0.10, None, 'timesteps/s')}
63         },
64         'large': {
65             'daint:gpu': {'perf': (2108.561, -0.10, None, 'timesteps/s')}
66         }
67     }
```

```
69 @run_after('init')
70 def setup_by_scale(self):
71     self.descr = f'LAMMPS GPU check (version: {self.scale})'
72     if self.scale == 'small':
73         self.valid_systems += ['dom:gpu']
74         self.num_tasks = 12
75         self.num_tasks_per_node = 2
76     else:
77         self.num_tasks = 32
78         self.num_tasks_per_node = 2
79
80     self.reference = self.refs_by_scale[self.scale]
```

Act upon the
different pipeline
stages

Make use of good software practices

```
83 @rfm.simple_test
84 class LAMPSCPUCheck(LAMMPSCheck):
85     valid_systems = ['daint:mc', 'eiger:mc', 'pilatus:mc']
86     refs_by_scale = {
87         'small': {
88             'dom:mc': {'perf': (4216.05, -0.10, None, 'timesteps/s')},
89             'daint:mc': {'perf': (2523.077, -0.10, None, 'timesteps/s')},
90             'eiger:mc': {'perf': (3807.095, -0.10, None, 'timesteps/s')},
91             'pilatus:mc': {'perf': (4828.986, -0.10, None, 'timesteps/s')}
92         },
93         'large': {
94             'daint:mc': {'perf': (2076.665, -0.10, None, 'timesteps/s')},
95             'eiger:mc': {'perf': (4922.81, -0.10, None, 'timesteps/s')},
96             'pilatus:mc': {'perf': (7247.484, -0.10, None, 'timesteps/s')}
97         }
98     }
```


ReFrame tests

What do they look like?

amber_check.py

Tests can
be based
on the
library of
tests

```
6 import contextlib
7 import reframe as rfm
8 from hpctestlib.sciapps.amber.nve import amber_nve_check
9
10
11 @rfm.simple_test
12 class cscs_amber_check(amber_nve_check):
13     modules = ['Amber']
14     valid_prog_environs = ['builtin']
15     extra_resources = {
16         'switches': {
17             'num_switches': 1
18         }
19     }
20     tags |= {'maintenance', 'production'}
21     maintainers = ['VH', 'SO']
22     num_nodes = parameter([1, 4, 6, 8, 16], loggable=True)
23     allref = {
24         1: {
25             'p100': {
26                 'Cellulose_production_NVE': (30.84, -0.10, None, 'ns/day'),
```

Skip tests if conditions are not met for a
given system

```
92 @run_after('setup')
93 def skip_if_no_topo(self):
94     proc = self.current_partition.processor
95     pname = self.current_partition.fullname
96     if not proc.info:
97         self.skip(f'no topology information found for partition {pname!r}')
98
99 @run_after('setup')
100 def set_num_tasks(self):
101     if self.variant == 'cuda':
102         self.num_tasks_per_node = 1
103     else:
104         proc = self.current_partition.processor
105         pname = self.current_partition.fullname
106         self.num_tasks_per_node = proc.num_cores
107
108     self.num_tasks = self.num_nodes * self.num_tasks_per_node
```

CPU topology detection can be used to set
the number of tasks a test may need

Tests can be written focusing already in strong scaling

How can you contribute?

Make PRs!

- Contribute tests to run at CSCS <https://github.com/eth-cscs/cscs-reframe-tests>
- Contribute to the library of tests at <https://github.com/reframe-hpc/reframe>
- Contribute to the framework at <https://github.com/reframe-hpc/reframe>
- In case of questions open issues at <https://github.com/reframe-hpc/reframe>

ReFrame



<https://reframe-hpc.readthedocs.io>



<https://github.com/reframe-hpc/reframe>



<https://reframe-slack.herokuapp.com/>



[@ReFrameHPC](https://twitter.com/ReFrameHPC)




We are a community
and we are here to
help!



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Thank you for your attention.

Questions?

ReFrame



<https://reframe-hpc.readthedocs.io>



<https://github.com/reframe-hpc/reframe>



<https://reframe-slack.herokuapp.com/>



[@ReFrameHPC](https://twitter.com/ReFrameHPC)

<https://github.com/eth-cscs/cscs-reframe-tests>

