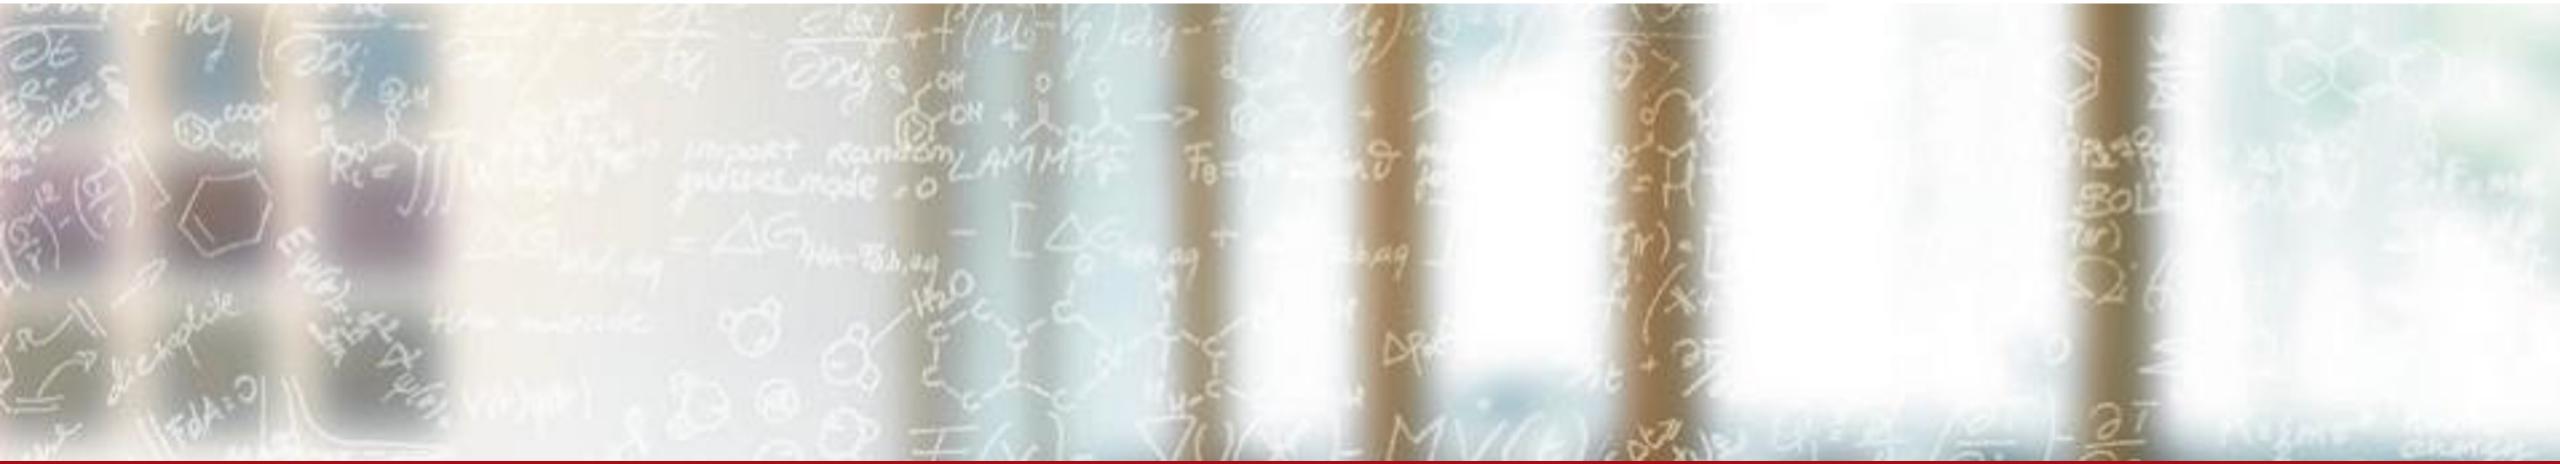




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH**zürich



# Spack: an efficient workflow in 30 minutes

CSCS User Lab Day 2022 – Meet the Swiss National Supercomputing Centre

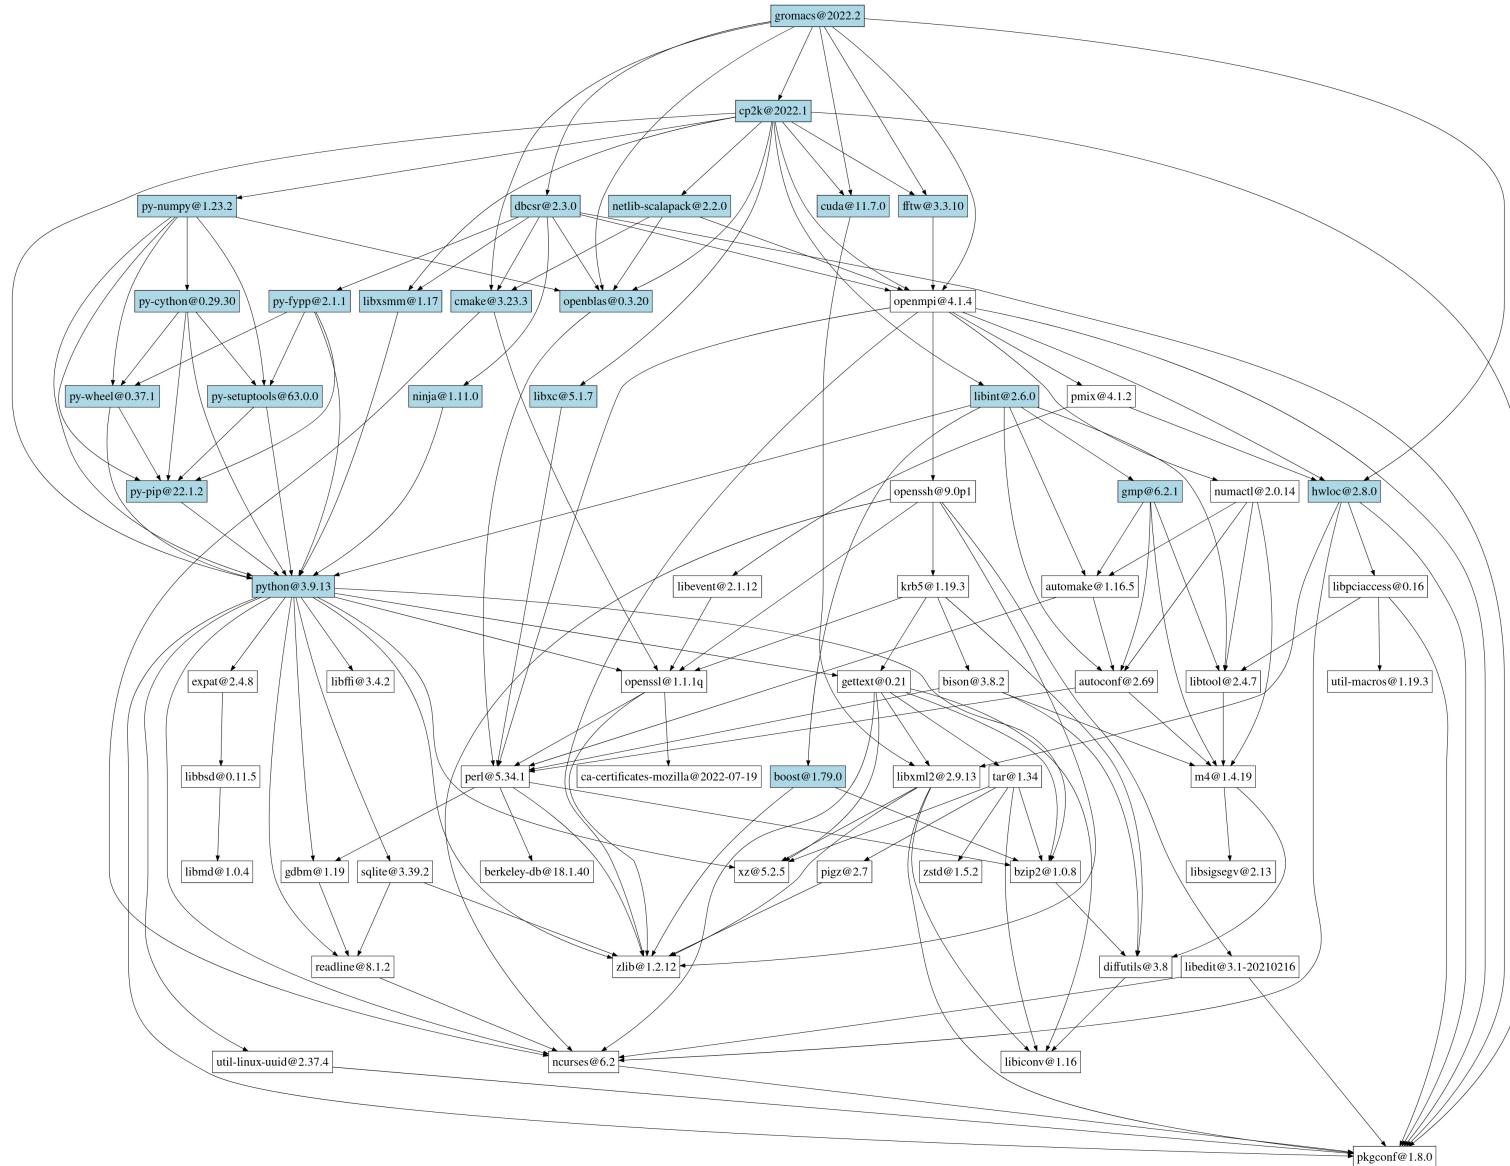
Harmen Stoppels and Alberto Invernizzi, CSCS

September 2nd, 2022

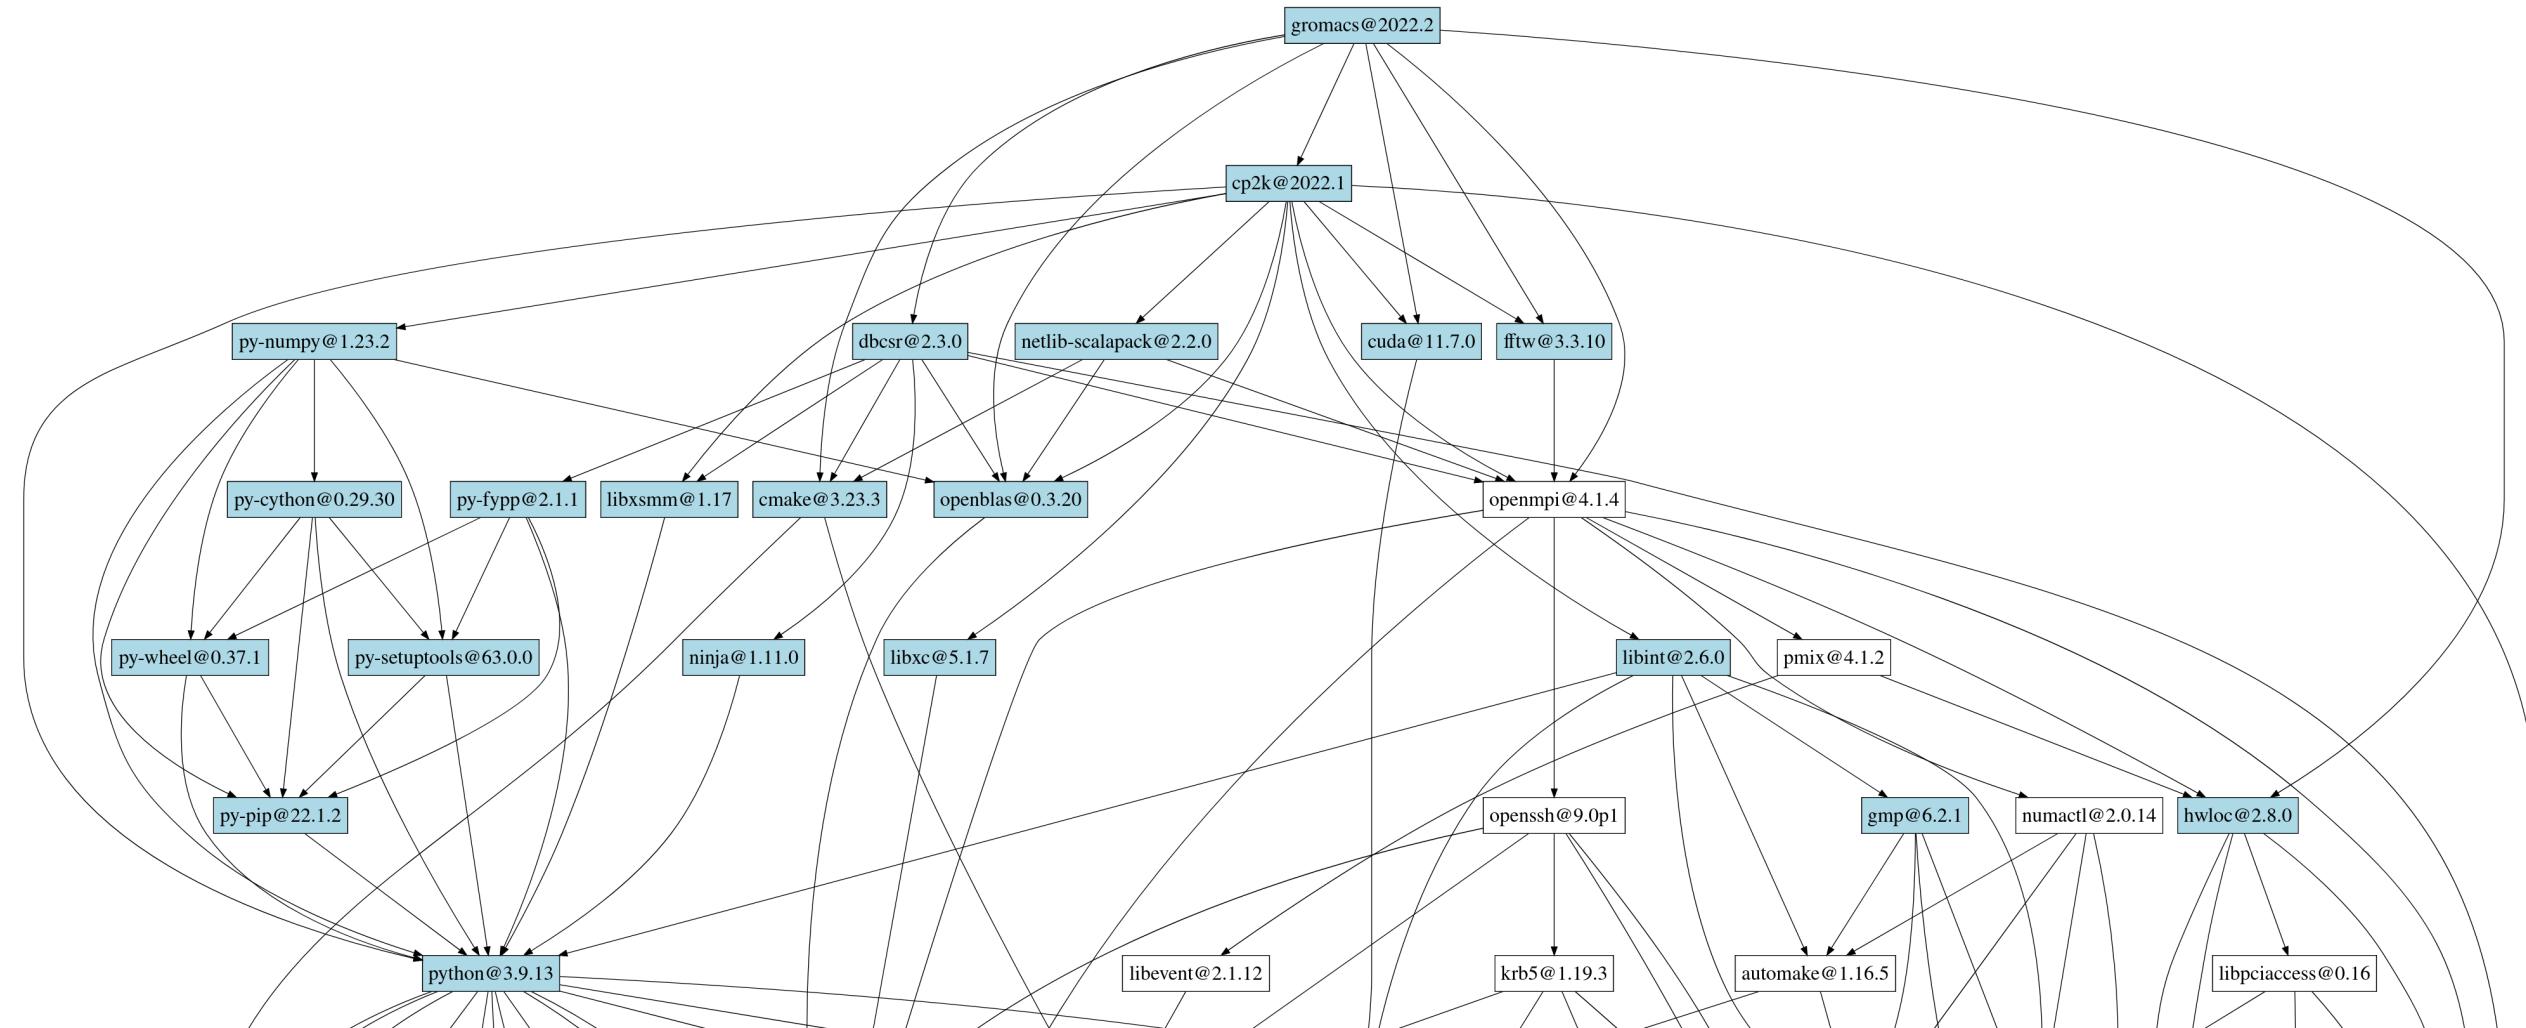
# Motivation

- Scientific software can have a complex set of dependencies
- The same version of a software package can come in many flavors
  - Enabling/disabling features
  - Different build types (debug vs release)
  - Different compilers (gcc, clang, ...)
- Manual installation is a tedious task

# Example: gromacs (with cp2k & cuda support)



# Example: gromacs (with cp2k & cuda support)



# How does Spack solve this issue?

- Spack is a package manager
- It resolves dependencies automatically
- It is customizable with a simple spec syntax



```
$ spack install gromacs +cp2k +cuda ^cp2k +cuda cuda_arch=70
```

The above does not install just gromacs, but also its dependencies, and dependencies of dependencies...



**CSCS**

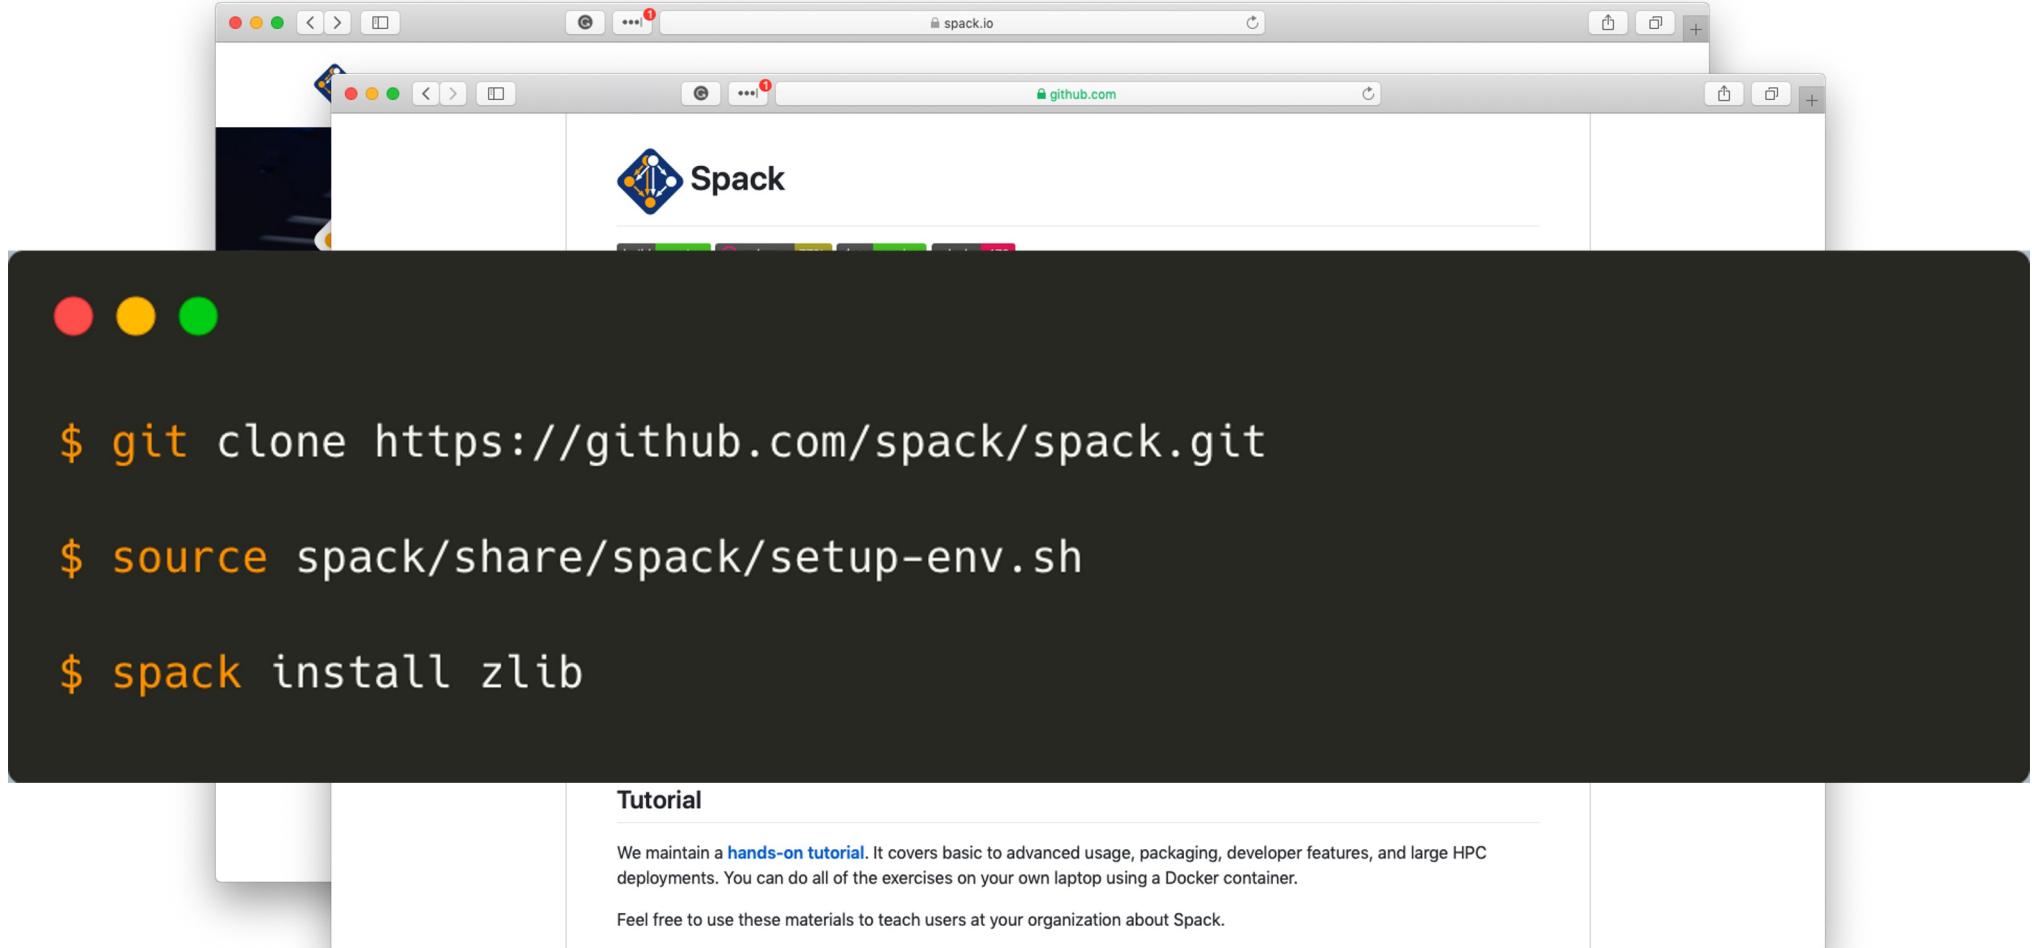
Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH**zürich

# Quick Start

---

# Let's start using it!



# SPEC...in Spack

*«[...] SPECS ARE MORE THAN A PACKAGE NAME, THEY CAN INCLUDE THE VERSION, ARCHITECTURE, COMPILE OPTIONS AND DEPENDENCIES»*

hpx@1.3.0 %gcc@8.4

- **@**
- **%**
- **gcc@4.7.3)**
- **+ or - or ~**
- **name=<value>**
- **name=<value>  
and ldlibs})**
- **target=<value> os=<value>**
- **^<dependency-spec>**

```
→ spack git:(develop) ✘ spack help --spec
spec expression syntax:

package [constraints] [^dependency [constraints] ...]

package
/hash                                              any package from 'spack list', or
                                                     unique prefix or full hash of
                                                     installed package

constraints:
  versions:
    @version                                         single version
    @min:max                                         version range (inclusive)
    @min:                                            version <min> or higher
    @:max                                             up to version <max> (inclusive)

  compilers:
    %compiler                                         build with <compiler>
    %compiler@version                                build with specific compiler version
    %compiler@min:max                               specific version range (see above)

  variants:
    +variant                                          enable <variant>
    -variant or ~variant                            disable <variant>
    variant=value                                    set non-boolean <variant> to <value>
    variant=value1,value2,value3                  set multi-value <variant> values

  architecture variants:
    platform=platform                             linux, darwin, cray, bgq, etc.
    os=operating_system                         specific <operating_system>
    target=target                                 specific <target> processor
    arch=platform-os-target                      shortcut for all three above

  cross-compiling:
    os=backend or os=be                           build for compute node (backend)
    os=frontend or os=fe                          build for login node (frontend)

  dependencies:
    ^dependency [constraints]                     specify constraints on dependencies
    ^/hash                                         build with a specific installed
                                                    dependency

examples:
  hdf5                                              any hdf5 configuration
  hdf5 @1.10.1                                       hdf5 version 1.10.1
  hdf5 @1.8:                                         hdf5 1.8 or higher
  hdf5 @1.8: %gcc                                    hdf5 1.8 or higher built with gcc
  hdf5 +mpi                                         hdf5 with mpi enabled
  hdf5 ~mpi                                         hdf5 with mpi disabled
  hdf5 +mpi ^mpich                                  hdf5 with mpi, using mpich
  hdf5 +mpi ^openmpi@1.7                            hdf5 with mpi, using openmpi 1.7
  boxlib dim=2                                      boxlib built for 2 dimensions
  libdwarf %intel ^libelf%gcc
    libdwarf, built with intel compiler, linked to libelf built with gcc
  mvapich2 %pgi fabrics=psm,mrall,sock
    mvapich2, built with pgi compiler, with support for multiple fabrics
```

TO SPECIFY THE COMPILER, COMPILER  
 » official docs (spack.readthedocs.io)

ne ~tools +examples

an optional compiler version (gcc or

clang, -qt, or ~qt) for boolean variants  
 selected to boolean variants  
 flags, cxxflags, fflags, cppflags, ldflags,

)

dependency on another package



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH**zürich

# A successful workflow: environments setup

---

# Creating an environment

1. Create an environment
2. Add a few packages or specs
3. Concretize
  - a. Resolves dependencies
  - b. Shows what will get installed
4. Install

# Creating an environment

1. `$ spack env create my_env`
2. `$ spack -e my_env add py-scipy ^python@3.10`
3. `$ spack -e my_env concretize`
4. `$ spack -e my_env install`

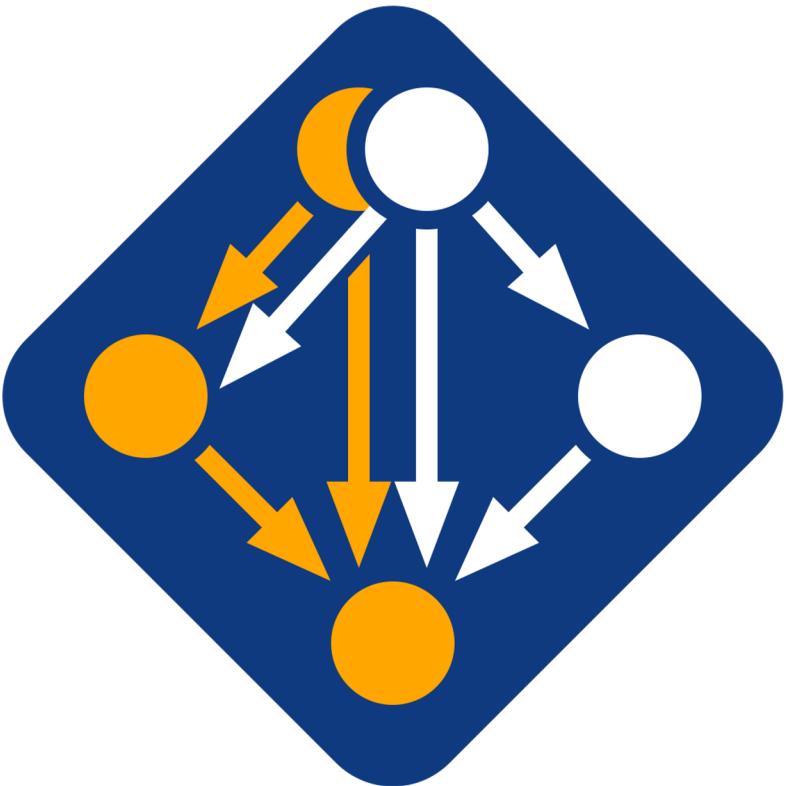
*NOTE: spack -e my\_env makes subcommands work within the specified environment*

```
[my_env] $ spack -e my_env concretize
- ldnsr6x py-scipy@1.8.1%gcc@11.1.0 arch=linux-ubuntu20.04-zen2
- n4b7nxt ^openblas@0.3.20%gcc@11.1.0~bignuma~consistent_fpcsr~ilp64+locking...
[+] 7eix2j5      ^perl@5.30.0%gcc@11.1.0~cpnm+shared+threads arch=linux-ubuntu...
- y3k6mww      ^py-cython@0.29.30%gcc@11.1.0 arch=linux-ubuntu20.04-zen2
- x37gqeq      ^py-pip@22.1.2%gcc@11.1.0 arch=linux-ubuntu20.04-zen2
- glduetc      ^python@3.10.6%gcc@11.1.0+bz2+ctypes+dbm~debug+libxml2+lzm...
[+] 7bdbbqq      ^bzip2@1.0.8%gcc@11.1.0~debug~pic+shared arch=linux-ub...
[+] d2y4dtt      ^diffutils@3.8%gcc@11.1.0 arch=linux-ubuntu20.04-z...
[+] eudajht      ^libiconv@1.16%gcc@11.1.0 libs=shared,static a...
[+] 355ji4i      ^expat@2.4.8%gcc@11.1.0+libbsd arch=linux-ubuntu20.04-...
[+] n3req45      ^libbsd@0.11.5%gcc@11.1.0 arch=linux-ubuntu20.04-z...
[+] vpv4x3u      ^libmd@1.0.4%gcc@11.1.0 arch=linux-ubuntu20.04...
[+] dvjhmhy      ^gdbm@1.19%gcc@11.1.0 arch=linux-ubuntu20.04-zen2
[+] sepfikp      ^readline@8.1.2%gcc@11.1.0 arch=linux-ubuntu20.04-...
[+] yv6danh      ^ncurses@6.2%gcc@11.1.0~symlinks+termib abi=n...
[+] qi3rmzy      ^pkgconf@1.8.0%gcc@11.1.0 arch=linux-ubunt...
[+] dgk5wrw      ^gettext@0.21%gcc@11.1.0+bzip2+curses+git~libunistring...
[+] s4imhr2      ^libiconv@1.16%gcc@11.1.0 libs=shared,static arch=...
[+] p7orov2      ^libxml2@2.10.1%gcc@11.1.0~python arch=linux-ubunt...
[+] l4oscwa      ^pkgconf@1.8.0%gcc@11.1.0 arch=linux-ubuntu20....
```

# What's in an environment?



```
spack:  
concretizer:  
    unify: true  
specs:  
- py-scipy  
- python@3.10
```



```
{
  "_meta": {
    "file-type": "spack-lockfile",
    "lockfile-version": 4,
    "specfile-version": 3
  },
  "roots": [
    {
      "hash": "ldnsr6xrkr3dfskktaawbep7y5rekcfak",
      "spec": "py-scipy"
    },
    {
      "hash": "glduetcfm24cxra64jy5lbtspok2z",
      "spec": "python@3.10"
    }
  ],
  "concrete_specs": {
    "ldnsr6xrkr3dfskktaawbep7y5rekcfak": {
      "name": "py-scipy",
      "version": "1.8.1",
      "dependencies": [
        {
          "spec": "numpy@1.19.2%gcc@9.2.0",
          "hash": "ldnsr6xrkr3dfskktaawbep7y5rekcfak"
        }
      ]
    }
  }
}
```

# Modifying an environment

Edit the **spack.yaml** file directly:

```
$ spack -e my_env config edit
```

Or

```
$ spack location -e my_env  
/path/to/my_env  
  
$ vim /path/to/my_env/spack.yaml
```

# Modifying an environment (different ways, same result)

```
spack:  
  concretizer:  
    unify: true  
  specs:  
    - py-scipy  
    - python@3.10  
    - intel-oneapi-mkl
```

```
spack:  
  specs:  
    - py-scipy  
    - intel-oneapi-mkl  
  concretizer:  
    unify: true  
  packages:  
    python:  
      require: '@3.10'
```

```
spack:  
  specs:  
    - py-scipy ^python@3.10 ^intel-oneapi-mkl
```

# Iterating concretization

1. 

```
$ spack -e my_env config edit
```
2. 

```
$ spack -e my_env concretize --force
```
3. 

```
$ spack -e my_env install
```



Iterate until happy!



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH**zürich

# A successful workflow: environments usage

---

# Using the installed environment

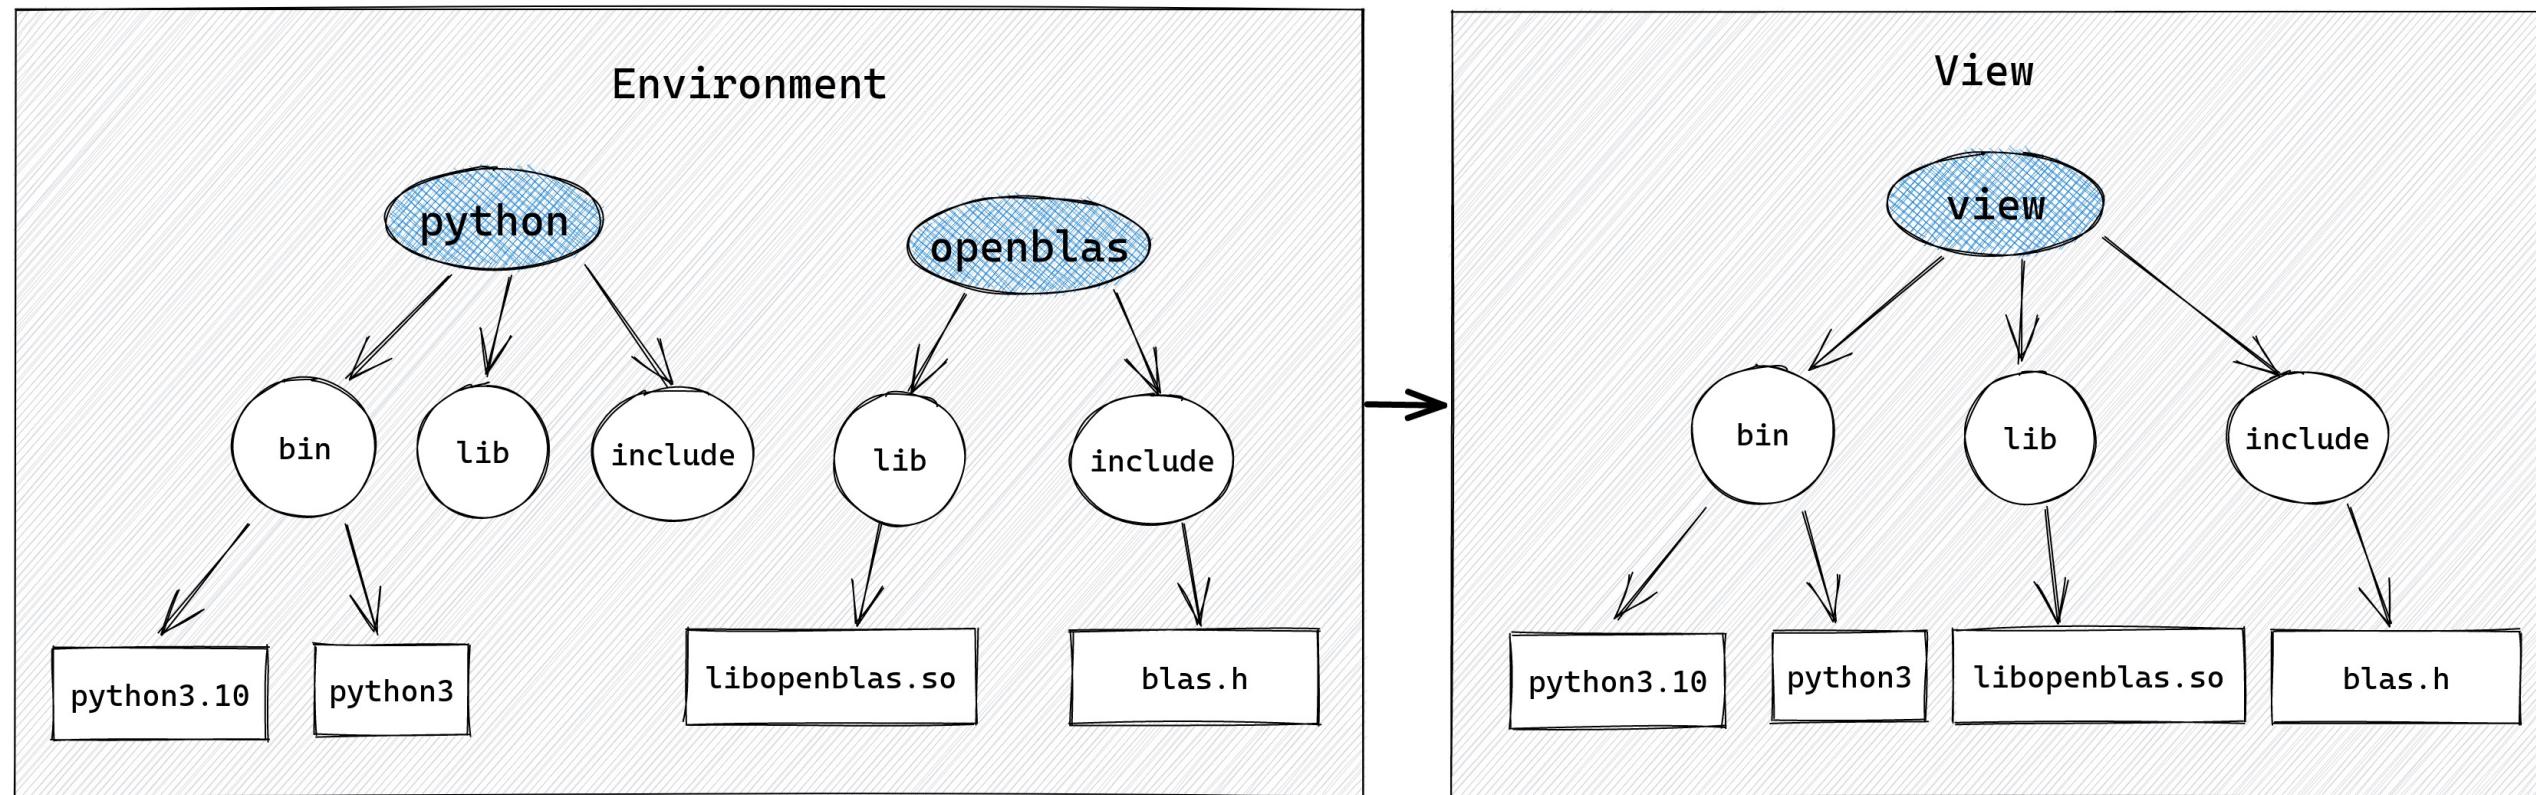
Typical interaction with a programming environment:

```
module swap PrgEnv-cray PrgEnv-gnu
module load gcc/10
module load ...
```

Not without issues

- **Consistency**: compatibility & dependencies across modules
- **Opaqueness**: many, many environment variables
- **Persistence**: what modules did I load last time?

# Using the installed environment: views



# Using the installed environment: views

- Persistent on disk
- Just a handful of environment variables
- Consistent set of libraries

```
$ srun ./view/bin/my_app
```

```
$ gcc -o app -I view/include app.c -L view/lib -lfoo
```

# Configuring views

```
spack:  
  view: /path/to/view  
  concretizer:  
    unify: true  
  specs:  
    - py-scipy  
    - python@3.10  
    - intel-oneapi-mkl
```



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH**zürich

# Spack at CSCS

---

# Spack at CSCS: Piz Daint

Many packages are already provided by the system, including optimized libraries (e.g. cray-mpich) and compilers.

CSCS provides spack configuration files for default CDT/CPEs that can be easily used to let your personal spack instance identify main packages, so that they can be used as dependencies in spack.

```
$ module load daint-gpu          # or module load daint-mc  
$ module load spack-config
```

# Spack at CSCS: Alps

- Many Spack packages come pre-installed
  - Spack 0.18+ prefers reuse over the latest the greatest
  - Decoupled from the Cray PE
- 
- As a user: less time spent configuring and building from sources

# Environments and replicability



spack.yaml

vs



spack.lock

describes **requirements**

- they can be easily share with users on different machines
- what is needed to get a correct working environment

describes **how to exactly reproduce the full installation:**

- it's safer to use this only on the same machine
- reproduce exactly the same working environment

```
$ spack env create new_env /path/to/spack.{yaml/lock}
```

you can use any file you want, or if you want to clone another environment

```
$(spack location -e other_env)/spack.yaml
```

# ...and it's not just for HPC!

It is a useful tool in your local development environment too!

- It keeps libraries organized with all the variants
  - debug/release
  - with or without a specific option
  - different versions
- It eases the installation process
  - and it builds from source with the microarchitecture optimizations

It is not just for libraries! You can find also tools like

- cmake
- tmux
- ...

And also compilers!

- gcc
- clang
- ...

# References



Documentation

<https://spack.readthedocs.io/>



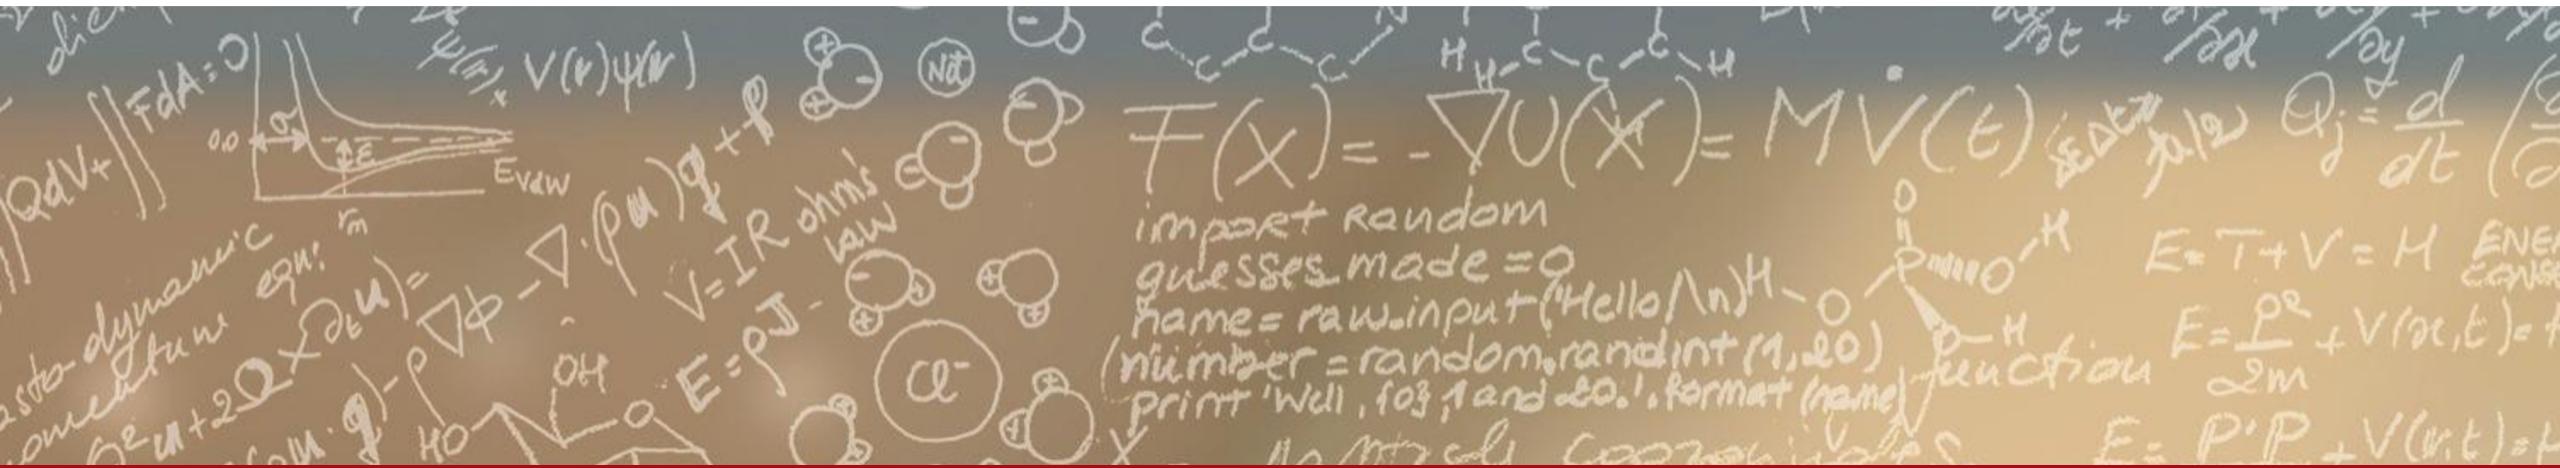
GitHub Repo

<https://github.com/spack/spack>



Slack channel

<https://slack.spack.io>



The end A new beginning

Thank you for your attention.