




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Building and Using HPC Software on Alps: CPE and uenv

CSCS User Day

Ben Cumming

September 2, 2024

# Today's Talk

I will present an overview of how to build and find HPC software on Alps

- Scientific software provided by CSCS on Alps
  - Cray Programming Environment
  - uenv
- uenv introduction
- uenv hands on
  - Getting started: exploring and managing uenv

We won't cover containers in this presentation, see:  
“ML and PyTorch in Containers” by Dr Nicholas Browning at 14:30 today

# Nomenclature

- Daint: the new Alps vCluster with GH200 nodes
- Daint-XC: “old daint”
- CPE: Cray Programming Environment

---

# Cray Programming Environment

# CPE: The Cray Programming Environment

HPE/Cray provide a programming environment that is familiar to all users:

- Daint XC - Cray Development Toolkit (CDT)
- Alps - Cray Programming Environment (CPE)

On Alps it is installed in `/opt/cray/`

- Delivered as a collection of RPMs - system engineers select the set of RPMs to install (e.g. do we install PrgEnv-Intel on AMD CPU nodes?)
  - CSCS can choose which RPMs to install, but we don't have much flexibility to modify the packages or their configuration
- Any change requires rebuilding the node image and rebooting nodes
- Vertically integrated
- Released on a 3 month cadence

Users configure the environment using modules

# Using the CPE

On Daint-XC the CPE is loaded by default with a default configuration.

On Alps you get a clean environment on login, and load CPE with a module

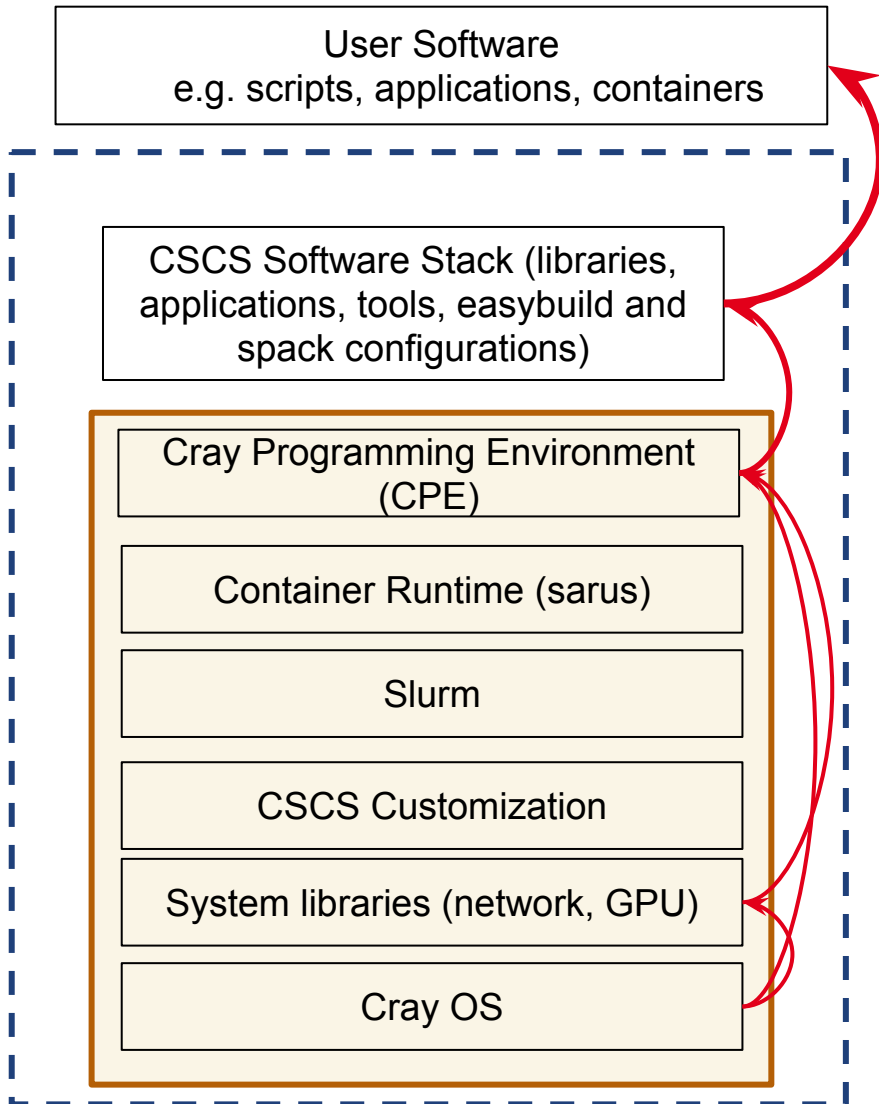
```
>> module list
No modules loaded
>> module avail
----- /etc/cscs-modules -----
  cray/23.12
>> module load cray
>> module list

Currently Loaded Modules:
 1) craype-arm-grace
 2) libfabric/1.15.2.0
 3) craype-network-ofi
 4) xpmem/2.8.2-1.0_5.1__g84a27a5.shasta
 5) PrgEnv-cray/8.5.0
 6) cce/17.0.0
 7) ...

>> module avail cray-mpich

---- /opt/cray/pe/lmod/modulefiles/comnet/crayclang/17.0/ofi/1.0 ----
  cray-mpich-abi/8.1.28    cray-mpich/8.1.28 (L)
```

# Providing software via CPE is challenging



The Cray Programming Environment is complex by necessity:

- Modules provide a combinatorial set of libraries and tools that serve as many use cases as possible on an increasing number of hardware types.
- Integration is provided by HPE: once an issue is identified HPE have to fix the issue in a future release
  - Long latency between issues reporting and the fix available on Daint.
- Each new release requires extensive testing to check that issues have been fixed
  - And to identify the inevitable new issues

**The Cray PE is the best configuration from any vendor in my experience. These challenges affect all HPC clusters.**

# Stability vs. Bug Fixes and New Features

**Any feedback that in your opinion can help us improve the HPC environment?**

*Would it be possible to keep older versions?*

*I am very satisfied with the HPC environment on Piz Daint*

*Compilers that support the newest C++ standards as well as possible*

*I need a stabler environment: older versions of the software tools disappear too quickly, which means I have to rebuild my stack every few months.*

*Please regularly update C++ and CUDA compilers*

By providing an environment on CPE it is very difficult to meet all requirements

- Regular updates are required to fix bugs, maintain security and provide updated versions of tools.
- The latest versions of compilers can't be installed before they are packaged by HPE and tested by CSCS.
- It is impractical to maintain:
  - Full stacks on top of more than one CPE
  - More than 2-3 CPE on a system





# CPE Support at CSCS

**CSCS will no longer provide software built using CPE for users**

CPE will still be provided on Daint

- It is provided on vClusters for users who have a hard dependency on Cray tools
  - contact me if you have concerns about dependencies on CPE
- If there is a bug or performance issue with CPE, CSCS will forward the HPE
  - this is a similar level of support to what CSCS provides for CPE in the past

**CSCS is focussing on [uenv](#) for deploying software on Alps vClusters**



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

---

**uenv**

# uenv are self-contained software stacks

uenv are built for a (**system**, **uarch**) pair.

There are 5 target microarchitecture (**uarch**):

- **zen2/zen3** CPU-only 2 x AMD Rome/Milan
- **a100** 4 NVIDIA A100 GPUs + 1 AMD CPU socket
- **mi200** 4 AMD Mi250x GPUs + 1 AMD CPU socket
- **gh200** 4 x **GH200**

Each **system** provides some key dependencies

- **libfabric** and **xpmem**: network libraries
- **slurm**
- customization points for the target audience

# uenv are self-contained software stacks

A uenv is two components:

## 1. A squashfs file

- *“a read-only file system that lets you compress whole file systems or single directories, write them to ordinary files, and then mount them using a loopback device.” [The Linux Documentation Project](#)*
- A single file that contains everything in a working environment

## 2. Meta data:

- information about the uenv build (when, where, who)
- the recipe that was used to build
- information about the contents of the uenv
- **environment configurations**

uenv have to be available on the local file system to be used

Store in repository, which is a directory with:

- A database: `index.db`
- A **hashed path** for each uenv that contains
  - the squashfs image `store.squashfs`
  - meta data: `env.json`

**index.db**

```
images/41fb...49c9/store.squashfs
images/41fb...49c9/meta/configure.json
images/41fb...49c9/meta/env.json
images/41fb...49c9/meta/recipe/
images/95fc...f8d6/meta/configure.json
images/95fc...f8d6/meta/env.json
images/95fc...f8d6/meta/recipe/
images/95fc...f8d6/store.squashfs
```

## uenv can be deployed at any time

uenv are a single file, that is mounted dynamically.

There is minimal coupling between each uenv and the underlying OS

1. use cases that need stability can continue to use the same uenv
  - in extreme cases where a low level upgrade breaks the uenv the same recipe can be rebuilt
2. updates and bug fixes can be deployed immediately without affecting old images or rebooting nodes
  - we can deploy the latest version of cray-mpich in a uenv within an hour of it being released

---

# Hands On

A picture is worth a thousand words

*and / or*

Watching somebody debug their own tool live is a great way to learn

- follow along in the terminal
- ... or watch the live demo
- ask questions!



---

# Documentation

# Uenv documentation

The most accessible docs are available on the command line

```
> uenv --help  
> uenv image --help  
> uenv image pull --help  
# help for the slurm plugin arguments is available via srun/sbatch:  
> srun --help | less
```

The CSCS knowledge base has more information:

<https://confluence.cscs.ch/x/bYDTKQ>

---

## Future work

## we are very busy improving the uenv experience and features

- user-managed uenv images in repos
- community-managed uenv deployment
  - separate repositories with uenv recipes and deployment pipelines
  - managed by communities, e.g. weather and climate
  - deployed directly to vClusters
- storage policies:
  - ensure that uenv are not affected by scratch deletion policies

There are corner cases that we are always discovering, e.g. running heterogenous slurm jobs:

```
srun -n8 -N2 --uenv=cp2k      ./wrap-cp2k.sh : \  
      -n4 -N1 --uenv=gromacs  ./wrap-gmx.sh
```

## please give your feedback

The uenv development team is aware of two types of uenv users:

- those who have simple problems and questions that are fixed quickly
- those who have complicated challenges that take significant effort to understand and fix

We don't have much information about two key groups

- happy users who don't need help or give feedback
- unhappy users who suffer in silence.

**Please give feedback about what works and what doesn't**

---

# Questions